

CIML Summer Institute Hands-On Exercises



Hands-On Exercises

8:00 - 8:45 -- Writing & Sharing Computational Analyses in Jupyter Notebooks

8:45 - 9:45 -- Spark

9:45 - 10:00 -- Break

10:00 - 10:30 -- Spark Continued

10:30 - 11:00 -- Deep Learning

11:00 - 12:00 -- Lunch

12:00 - 1:00 -- Deep Learning Transfer Learning

1:00 - 1:45 -- Deep Learning Time Series & Others

1:45 - 2:00 -- Wrap-Up

Spark Hands-On

Mai H. Nguyen, Ph.D.

SPARK PROGRAM STRUCTURE

- **Start Spark session**
 - `spark = SparkSession.builder.config(conf=conf).getOrCreate()`
- **Create distributed dataset**
 - `df = spark.read.csv("data.csv",header="True")`
- **Apply transformations**
 - `new_df = df.filter(col("dept") == "Sales")`
- **Perform actions**
 - `df.collect()`
- **Stop Spark session**
 - `spark.stop()`


START SPARK SESSION

```
import pyspark
from pyspark.sql import SparkSession
```

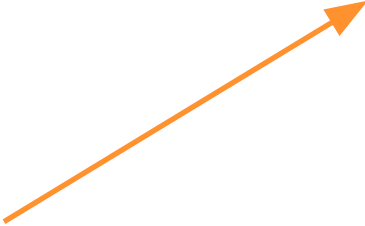
```
conf = pyspark.SparkConf().setAll([
    ('spark.master', 'local[*]'),
    ('spark.app.name', 'PySpark Demo')])
```

```
spark = SparkSession.builder.config(conf=conf).getOrCreate()
```


Use * to use all available cores, or integer value to specify number of cores to use



Configuration parameters for Spark session



Get existing Spark session or create new one



LOAD DATA

- Loading data from local file system

```
df = spark.read.text("file:///<path>/<file>.txt")
```

```
df = spark.read.csv("file:///<path>/<file>.csv",  
                    header=True).cache()
```

- Loading data from HDFS

```
df = spark.read.text("hdfs:///<path>/<file>.txt") \  
    .cache()
```

```
df = spark.read.csv("hdfs:///<path>/<file>.csv",  
                    header=True).cache()
```

Indicates whether
column headers exist



Cache data



CHAINING

```
text_file = sc.textFile("hdfs://...")

counts = text_file.flatMap(lambda line: line.split(" ")) \
    .map(lambda word: (word, 1)) \
    .reduceByKey(lambda a, b: a + b)

counts.saveAsTextFile("hdfs://...")
```

**Chaining: Making multiple
method calls on same object**



**Line continuation
indicator**



RDD Wordcount

DROP ROWS WITH NULLS

- **Drop rows with null values**

```
df.dropna()  
df.dropna(how='any')  
df.dropna(how='all')
```

- **Check number of rows before and after dropping rows**

```
df.count()
```


FILL IN MISSING VALUES

- **Replace null values with empty string**

```
df.na.fill('')
```

- **Count number of rows with nulls before and after filling nulls**

```
df.count()
```

PARTITION DATA

- Partition available data into train and test data sets

```
train_df, test_df = df.randomSplit(0.8, 0.2), seed=<seed>)
```

Percentage of
samples for train
dataset



Percentage of
samples for test
dataset



CREATE FEATURE VECTOR COLUMN

- **Create feature vector column**
 - Combines given list of columns into single vector column
 - To feed data to machine learning models

```
from pyspark.ml.feature import VectorAssembler
```

```
features = ['air_temp', 'relative_humidity']  
assembler = VectorAssembler(inputCols=features,  
                             outputCol='featureVector')
```

```
features_df = assembler.transform(df)  
features_df.show()
```

air_temp	relative_humidity
62.96	63.9



air_temp	relative_humidity	featureVector
62.96	63.9	[62.96, 63.9]

New column
appended to
features_df

SCALE DATA

- **Scale input data values**

- Standardize values to have zero mean and unit standard deviation
- Each feature is scaled separately
- Create scale transformer using train data, then apply to train/test data

```
from pyspark.ml.feature import StandardScaler
scaler = StandardScaler(inputCol="features_unscaled",
                        outputCol="features_scaled",
                        withStd=True, withMean=True)
scalerModel = scaler.fit(train_df)
```

```
train_df = scalerModel.transform(train_df)
test_df  = scalerModel.transform(test_df)
```

BUILD MODEL


- **Build decision tree classifier**

- Create model
- Use fit() to train model

```
from pyspark.ml.classification import DecisionTreeClassifier
```

```
dt = DecisionTreeClassifier(  
    featuresCol='featureVector',  
    labelCol='label',  
    predictionCol='prediction')
```

Can specify name of
columns for features,
labels, and predictions



```
dt_model = dt.fit(<train>)
```

APPLY MODEL

- **Apply trained model**
 - Use transform()

```
predictions = <model>.transform(<data>)
```

Will have <data> and
new 'prediction'
column appended at
the end



Trained
model



Input data



EVALUATE CLASSIFICATION MODEL

- **Evaluator for classification model**
 - Calculates F1, precision, recall, accuracy

```
from pyspark.ml.evaluation import MulticlassClassificationEvaluator

mc_evaluator = MulticlassClassificationEvaluator(
    predictionCol='prediction',
    labelCol='label',
    metricName='f1')

mc_evaluator.evaluate(<predictions>)
```

Column with predictions

Column with labels

Evaluation metric

Contains predictions and labels

PySpark Cluster Analysis Hands-On

- **Data**
 - Weather station measurements
- **Task**
 - Perform cluster analysis to identify different weather patterns
- **Approach**
 - Spark k-means
- **Notebook**
 - `pyspark-clustering.ipynb`


Dataset Description

- **Measurements from weather station on Mt. Woodson, San Diego**
- **Air temperature, humidity, wind speed, wind direction, etc.**
- **Three years of data: Sep. 2011 - Sep. 2014**
 - minute_weather.csv: measurement every minute
- **Source**
 - <http://hpwren.ucsd.edu>

A topographic map of San Diego and the surrounding regions, including parts of Imperial, Riverside, and Orange counties. The map shows various mountain peaks and geographical features. Labeled peaks include Toro Peak, Sky Oaks, Warner Springs, Mesa Grande, Big Black Mt., Santa Ysabel, Mt. Woodson, West Ramona, Mt. Laguna, Lyons Peak, and Red Mt. Other labeled locations include Wildomar, Murrieta, Fallbrook, Boucher Hill, Puerta La Cruz, Valley Center, Escondido, Carlsbad, Vista, San Marcos, Encinitas, Poway, Santee, El Cajon, Lemon Grove, La Presa, Chula Vista, San Diego, Oceanside, and San Clemente Island. The map also shows the Cahuilla and San Jacinto Reservations, the Cleveland National Forest, and the Anza-Borrego Desert State Park. A red arrow points from the bottom left towards the center of the map, and a black arrow points from the left towards San Clemente Island.

Cluster Analysis

- **Cluster analysis divides data into groups**
 - Grouping is based on some similarity measure.
 - Samples within a cluster are more similar to each other than to samples in other clusters.



<http://www-users.cs.umn.edu/~kumar/dmbook/index.php>

k-Means Clustering

- **Partitional**
 - Clusters are divided into non-overlapping subsets
- **Centroid-Based**
 - Cluster represented by central vector
- **Simple, classic clustering technique**
 - Data points are grouped into k clusters
 - Cluster defined by cluster mean

- **Algorithm**


Select k initial *centroids* (cluster centers)

Repeat


Assign each sample to closest centroid

Calculate mean of cluster to determine new centroid


Until some stopping criterion is reached




k-Means Clustering Illustration




Original samples




Initial Centroids




Assign Samples



Re-calculate Centroids




Assign Samples



Re-calculate
Centroids

Choosing Number of Clusters (k)

- **Elbow method**
 - Plot cluster evaluation metric (e.g., WSSE) vs. different values for k
 - “Elbow” in plot suggests value(s) for k



<http://stackoverflow.com/questions/6645895/calculating-the-percentage-of-variance-measure-for-k-means>

Evaluating Clustering Results

- **Within-Cluster Sum of Squared Error (WSSE)**
- For each sample, error is distance to centroid.
Then, **WSSE** is computed as:

$$WSSE = \sum_{i=1}^K \sum_{x \in C_i} \|x - m_i\|^2$$

x : data sample in cluster C_i

m_i : cluster centroid (i.e., mean of cluster)


$\|x - m_i\|^2$: Euclidean distance between m_i and x

Clustering Hands-On Overview

- **Setup**
 - Start Spark
 - Load modules
- **Load data**
 - Set schema
 - Read in data from “minute_weather.csv”
- **Explore data**
 - Look at schema, number of rows, summary statistics
- **Prepare data**
 - Drop nulls
 - Create feature vector
- **Perform k-means cluster analysis**
 - Use elbow plot to determine k
- **Evaluate clusters**
 - Plot cluster profiles
- **Stop Spark session**

Cluster Profile: Parallel Plots

```
utils.parallel_plot(centersNamed[centersNamed['relative_humidity'] < -0.5],  
numClusters, colors=colors_used);
```



Setup

- **Login to Expanse**

- Open terminal window on local machine
- `ssh login.expanse.sdsc.edu`

- **Pull latest from repo**

- `git pull`
- URL:

<https://github.com/ciml-org/ciml-summer-institute-2021>

Server Setup for PySpark - Command Line

- **In terminal window**
 - `start_spark`
 - Alias for: `galileo.sh launch --account sds184 --reservation 'ciml-day3' --partition 'shared' --ntasks-per-node 1 --cpus-per-task 4 --memory-per-node 8 --time-limit 03:00:00 --jupyter lab --notebook-dir "/home/${USER}" --sif /cm/shared/apps/containers/singularity/ciml/2021/pyspark-latest.sif --env-modules 'singularitypro' --bind /expanse,/scratch,/cvmfs --quiet`
- **To check queue**
 - `squeue -u $USER`

Data Setup

- In terminal window in Jupyter Lab, do the following:
- **Go to Spark directory**
`cd <CIML2021_dir>`
- **Link to data directory**
In -s /exppanse/lustre/projects/sds184/mhnguyen/ciml2021_data data

ls data
BookReviews_1M.txt BookReviews_5M.txt catsVsDogs daily_weather.csv
minute_weather.csv winequality-white.csv

PySpark Cluster Analysis Hands-On

- **Code**

- **pyspark-cluster.ipynb**
 - Notebook for hands-on
 - Replace <<FILL-IN>> with code
- **pyspark-cluster-w-outputs.ipynb**
 - Has cell outputs
- **utils.py**
 - Has utility functions

- **Resources**

- [Apache Spark™ - Unified Analytics Engine for Big Data](#)
- [PySpark Documentation — PySpark 3.1.2 documentation](#)
- [Spark SQL and DataFrames - Spark 3.1.2 Documentation](#)
- Python for Data Science Cheat Sheet (pdf)

PySpark Scaling Hands-On

- **Data**
 - BookReviews_5M.txt
 - Reference : <https://jmcauley.ucsd.edu/data/amazon/>
- **Notebook**
 - pyspark_demo.ipynb
- **To do**
 - Change number of cores: 1, 2, 4
 - Note difference in execution times


START SPARK SESSION

```
import pyspark
from pyspark.sql import SparkSession
```

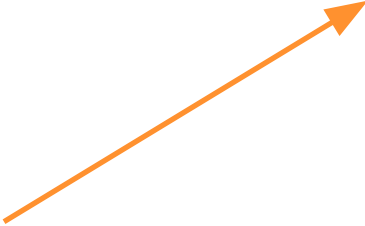
```
conf = pyspark.SparkConf().setAll([
    ('spark.master', 'local[*]'),
    ('spark.app.name', 'PySpark Demo')])
```

```
spark = SparkSession.builder.config(conf=conf).getOrCreate()
```

Use * to use all available cores, or integer value to specify number of cores to use



Configuration parameters for Spark session



Get existing Spark session or create new one




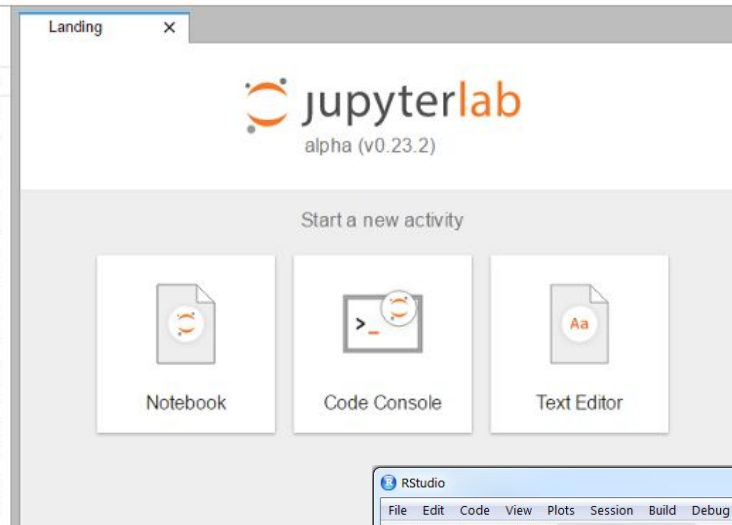

SparkR

- **R package that provides frontend to use Spark from R**
- **Supports distributed machine learning using R API**
- **Allows R script to connect to Spark cluster**
- **Can use with R shell or RStudio or other R IDEs.**

SparkR

- **Uses MLlib for machine learning functionality**
- **Familiar R syntax:**
 - Read contents of file into a Spark dataframe
 - `newdata <- read.df ("data.txt", source="csv")`
 - R formula operators for model fitting:
 - `model <- spark.randomForest(training, label ~ features, "classification", numTrees = 10)`
 - Get summary of fitted model
 - `summary(model)`
 - Apply model to make predictions
 - `predictions <- predict(model, testDF)`
 - Save model
 - `write.ml (model, "mymodel")`

Running SparkR Code





sparklyr



- **R interface for Spark**
- **R Legacy**
 - From Rstudio
 - Available on CRAN
- **Spark backend to dplyr and SQL**
 - Interactively manipulate Spark data using dplyr and SQL
- **Access to Spark functionality**
 - Interface to Spark MLlib algorithms
- **Connect to Spark clusters**

sparklyr

- **Setup**

- `install.packages("sparklyr")`
- `library(sparklyr)`
- `spark_install ()`

- **Connect to Spark**

- `sc <- spark_connect (master="local")`

- **Using dplyr**

- `flights_sdf %>% group_by(tailnum)`
`%>% filter(count > 20)`

- **Machine Learning**

- `model <- ml_random_forest (`
`train_sdf, quality ~ ., type="classification")`

sparklyr Functions

- **Spark Operations**
 - Manage Spark connections (e.g., `spark_config()`)
- **Spark Data Manipulation**
 - Read data in Spark DataFrame and perform operations (e.g., `spark_read_csv()`, `tbl_cache()`)
- **Feature Transformers**
 - Transform data (e.g., `ml_pca()`, `ft_bucketizer()`)
- **Distributed Machine Learning**
 - Access Spark Mllib algorithms (e.g., `ml_kmeans()`)
- **Streaming**
 - Support streaming data operations (e.g., `stream_read_json()`)
- **Extensions**
 - Interface to platforms for big data analysis, graph analytics, production

Machine Learning Algorithms in SparkR

Machine Learning

Algorithms

SparkR supports the following machine learning algorithms currently:

Classification

- `spark.logit`: Logistic Regression
- `spark.mlp`: Multilayer Perceptron (MLP)
- `spark.naiveBayes`: Naive Bayes
- `spark.svmLinear`: Linear Support Vector Machine

Regression

- `spark.survreg`: Accelerated Failure Time (AFT) Survival Model
- `spark.glm` or `glm`: Generalized Linear Model (GLM)
- `spark.isoreg`: Isotonic Regression

Tree

- `spark.gbt`: Gradient Boosted Trees for Regression and Classification
- `spark.randomForest`: Random Forest for Regression and Classification

Clustering

- `spark.bisectingKmeans`: Bisecting k-means
- `spark.gaussianMixture`: Gaussian Mixture Model (GMM)
- `spark.kmeans`: K-Means
- `spark.lda`: Latent Dirichlet Allocation (LDA)

Collaborative Filtering

- `spark.als`: Alternating Least Squares (ALS)

Frequent Pattern Mining

- `spark.fpGrowth`: FP-growth

Statistics

- `spark.kstest`: Kolmogorov-Smirnov Test

Spark Hands-On Topics

- Spark Review
- PySpark Examples
- PySpark Hands-On Exercise
- System Setup
- **SparkR**

Resources

- **Spark**
 - <https://spark.apache.org/>
- **PySpark API**
 - <https://spark.apache.org/docs/latest/api/python/index.html>
- **Spark DataFrame**
 - <https://spark.apache.org/docs/latest/sql-programming-guide.html>
!
- **MLlib**
 - <https://spark.apache.org/mllib/>
- **SparkR**
 - <https://spark.apache.org/docs/latest/sparkr.html>
- **SparkR API**
 - <https://spark.apache.org/docs/latest/api/R/>
- **sparklyr**
 - <https://spark.rstudio.com/>