

CIML Summer Institute Hands-On Exercises



Hands-On Exercises

8:00 - 8:45 -- Writing & Sharing Computational Analyses in Jupyter Notebooks

8:45 - 9:45 -- Spark

9:45 - 10:00 -- Break

10:00 - 10:30 -- Spark

10:30 - 11:00 -- MNIST & TensorBoard

11:00 - 12:00 -- Lunch

12:00 - 1:00 -- Deep Learning Transfer Learning

1:00 - 1:45 -- Deep Sequence Learning

1:45 - 2:00 -- Q&A

Deep Learning Transfer Learning Hands-On

Mai H. Nguyen, Ph.D.

WHAT IS TRANSFER LEARNING?

- **To overcome challenges of training model from scratch:**
 - Insufficient data
 - Very long training time
- **Use pre-trained model**
 - Trained on another dataset
 - This serves as starting point for model
 - Then train model on current dataset for current task

TRANSFER LEARNING APPROACHES

- **Feature extraction**

- Remove last fully connected layer from pre-trained model
- Treat rest of network as feature extractor
- Use features to train new classifier (“top model”)

- **Fine tuning**

- Tune weights in some layers of original model (along with weights of top model)
- Train model for current task using new dataset

CNNs FOR TRANSFER LEARNING

- **Popular architectures**
 - AlexNet
 - GoogLeNet
 - VGGNet
 - ResNet
- **All winners of ILSVRC**
 - ImageNet Large Scale Visual Recognition Challenge
 - Annual competition on vision tasks on ImageNet data

ImageNet

- **Database**

- Developed for computer vision research
- > 14,000,000 images hand-annotated
- > 22,000 categories

- **ILSVRC History**

- Started in 2010
- Image classification task: 1,000 object categories
- Image classification error rate
 - 2011: ~25% (conventional image processing techniques)
 - 2012: 15.3% (AlexNet)
 - 2015: 3.57% (ResNet; better than human performance)
 - 2016: 2.99% (16.7% error reduction)
 - 2017: 2.25% (23.3% error reduction)

TRANSFER LEARNING


Input



Output




*Learned
hierarchy*




Lee et al. 'Convolutional Deep Belief Networks for Scalable
Unsupervised Learning of Hierarchical Representations' ICML 2009

PRE-TRAINED MODEL




<https://www.cs.toronto.edu/~frossard/post/vgg16/>

TRANSFER LEARNING - FEATURE EXTRACTION



<https://www.cs.toronto.edu/~frossard/post/vgg16/>

TRANSFER LEARNING - FINE TUNING



<https://www.cs.toronto.edu/~frossard/post/vgg16/>

TRANSFER LEARNING EXERCISES

- **Data**
 - Cats and dogs images from Kaggle
- **Exercises**
 - Feature extraction
 - Use pre-trained CNN to extract features from images
 - Train neural network to classify cats/dogs using extract features
 - Fine tune
 - Adjust weights of last few layers of pre-trained CNN through training

DATA

- **Subset of Kaggle cats and dogs dataset**
- **Train**
 - 1000 cats + 1000 dogs
- **Validation**
 - 400 cats + 400 dogs
- **Test**
 - 400 cats + 400 dogs



TRANSFER LEARNING - FEATURE EXTRACTION


- **Data**

- Cats and dogs images from Kaggle

- **Method**

- Use VGG16 trained on ImageNet data as pre-trained model. Remove last fully connected layer.
- Extract features from pre-trained model and save
- Neural network then trained on extracted features to classify cats vs. dogs

TRANSFER LEARNING - FEATURE EXTRACTION




Source: <https://www.cs.toronto.edu/~frossard/post/vgg16/>

Feature Extraction Overview

- **Data**
 - Set image dimensions & location
 - Use ImageDataGenerator to read images from folder
- **Model**
 - Load model pre-trained on ImageNet data
 - Freeze weights in pre-trained model to use as feature extractor
 - Add top model to classify cats vs dogs
 - Model = Pre-trained base model + top model classifier
- **Train model**
 - Use training data to adjust model weights
 - Use validation data to determine when to stop training
- **Evaluate model**
 - Calculate accuracy, etc.
 - Perform inference on test images

Early Stopping

Using validation data to determine when to stop training to avoid overfitting



TRANSFER LEARNING - FINE TUNING


- **Data**

- Cats and dogs images from Kaggle

- **Method**

- Use VGG16 trained on ImageNet data as pre-trained model.
- Replace last fully connected layer with neural network trained from Feature Extraction hands-on.
- Fine tune last convolution block and fully connected layer.

TRANSFER LEARNING - FINE TUNING



Source: <https://www.cs.toronto.edu/~frossard/post/vgg16/>

WHEN & HOW TO FINE TUNE

- **New dataset is small & similar to original dataset**
 - Extract features from higher layer and feed to separate classifier
- **New dataset is large & similar to original dataset**
 - Fine tune top or all layers
- **New dataset is small & different from original dataset**
 - Extract features from lower layer and feed to separate classifier
- **New dataset is large & different from original dataset**
 - Fine tune top or all layers

<http://cs231n.github.io/transfer-learning/>

OTHER PRACTICAL TIPS

- **Learning rate**
 - Use very small learning rate for fine tuning. Don't want to destroy what was already learned.
- **Start with properly trained weights**
 - Train top-level classifier first, then fine tune lower layers.
 - Top model with random weights may have negative effects on when fine tuning weights in pre-trained model
- **Data augmentation**
 - Simple ways to slightly alter images
 - Horizontal/vertical flips, random crops, translations, rotations, etc.
 - Use to artificially expand your dataset

Fine Tune Overview

- **Data**
 - Set image dimensions & location
 - Use ImageDataGenerator to read images from folder
- **Model**
 - Load trained model from feature extraction code
 - Freeze weights up to last convolutional block
 - Weights in last convolutional block and top classifier will be adjusted during training
- **Train model**
 - Use training data to adjust model weights
 - Use validation data to determine when to stop training
- **Evaluate model**
 - Calculate accuracy, etc.
 - Perform inference on test images

Code

- **features_extract_tf.ipynb**
 - Transfer learning with feature extraction
- **finetune_tf.ipynb**
 - Transfer learning with fine tuning
- **Note**
 - Close features_extract_tf.ipynb before running finetune_tf.ipynb to avoid out-of-memory errors

Setup

- **Login to Expanse**

- Open terminal window on local machine
- `ssh login.expanse.sdsc.edu`

- **Pull latest from repo**

- `git pull`
- URL:

<https://github.com/ciml-org/ciml-summer-institute-2021>

Server Setup for TensorFlow - Portal

- **Expanse Portal**
 - <https://portal.expanse.sdsc.edu>
- **Parameters**
 - Account: sds184
 - Time limit (min): 180
 - Number of cores: 10
 - Memory required per node: 93 GB
 - GPUs: 1
 - Singularity image:
/cm/shared/apps/containers/singularity/ciml/2021/tensorflow-lat
est.sif
 - Environment module: singularitypro
 - Reservation: ciml-day3
 - Working directory: home
 - Type: JupyterLab

Server Setup for TensorFlow - Command Line

- **In terminal window**
 - `start_gpu`
 - Alias for: `galileo.sh launch --account 'sds184' --reservation 'ciml-day3' --partition 'gpu-shared' --cpus-per-task 10 --memory-per-node 93 --gpus 1 --time-limit 03:00:00 --jupyter 'lab' --notebook-dir "/home/${USER}" --env-modules 'singularitypro' --sif '/cm/shared/apps/containers/singularity/tensorflow/tensorflow-latest.sif' --bind '/expance,/scratch,/cvmfs' --nv --quiet`
- **To check queue**
 - `squeue -u $USER`

REFERENCES

- **TensorFlow tutorial**
 - https://github.com/tensorflow/docs/blob/master/site/en/tutorials/images/transfer_learning.ipynb
- **TensorFlow/Keras API**
 - https://www.tensorflow.org/api_docs/python/tf/keras/Model
- **Transfer Learning**
 - <http://cs231n.github.io/transfer-learning/>